

Graph and geometric deep learning

Presenters: Simone Scardapane, Indro Spinelli



SAPIENZA
UNIVERSITÀ DI ROMA



intelligent signal processing
and multimedia lab

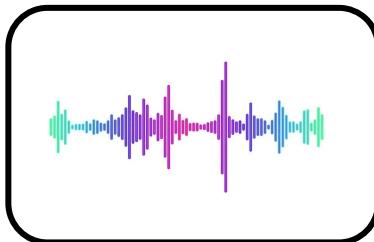
Enel Third Global Data Meet Up, April 22nd 2021

Introduction

On the importance of **graphs** in
deep learning

Data ingestion in deep learning

Audio



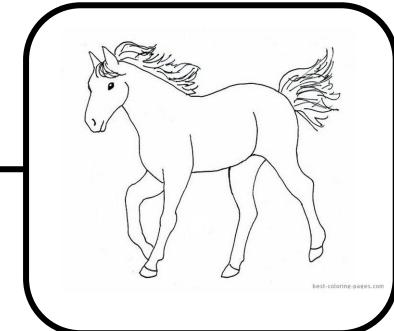
WaveNet, Wav2Vec, ...

Texts

Dies ist ein Blindtext. An ihm lässt sich vieles über die Schrift ablesen, in der er gesetzt ist. Auf den ersten Blick wird der Grauwert der Schriftfläche sichtbar. Dann kann man prüfen, wie gut die Schrift zu lesen ist und wie sie auf den Leser wirkt. Dies ist ein Blindtext. An ihm lässt sich vieles über die Schrift ablesen, in der er gesetzt ist. Auf den ersten Blick wird der Grauwert der Schriftfläche sichtbar. Dann kann man prüfen, wie gut die Schrift zu lesen ist und wie sie auf den Leser wirkt.

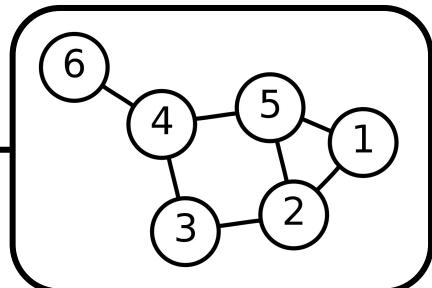
Word embeddings,
Transformers, ...

Images



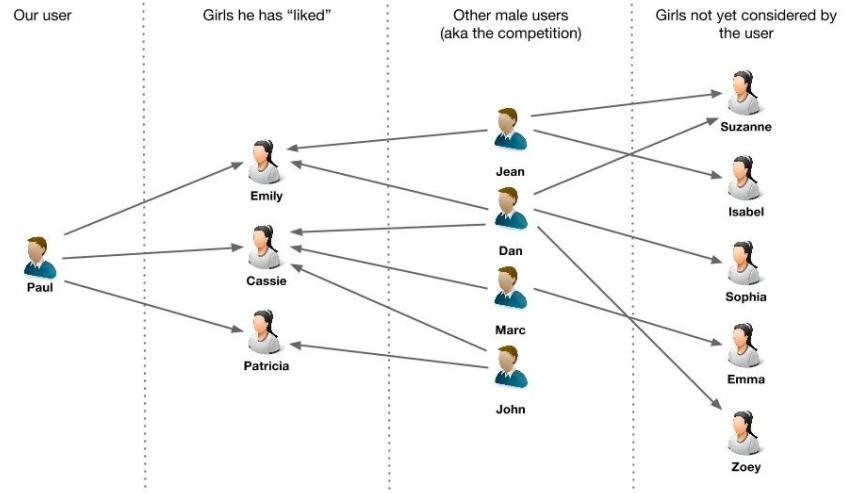
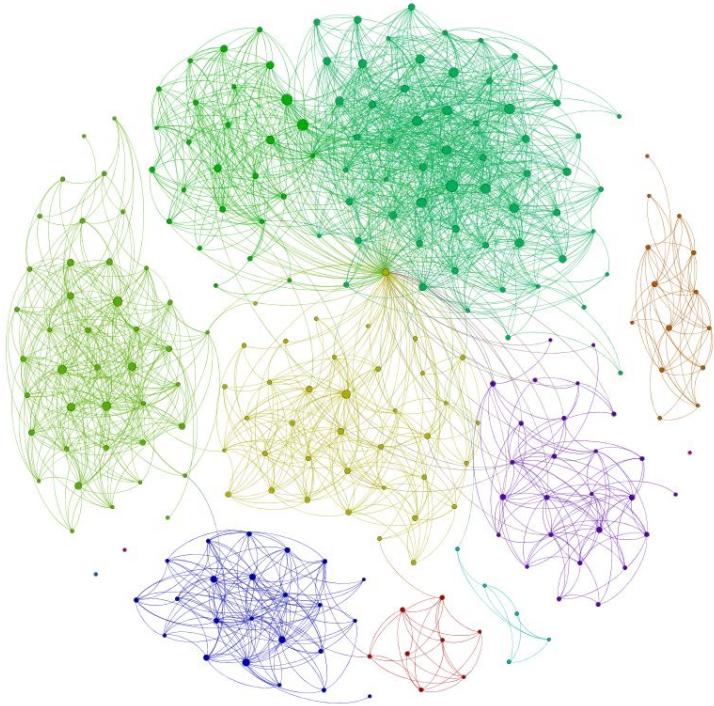
CNNs, Vision
Transformers, ...

Graphs



???

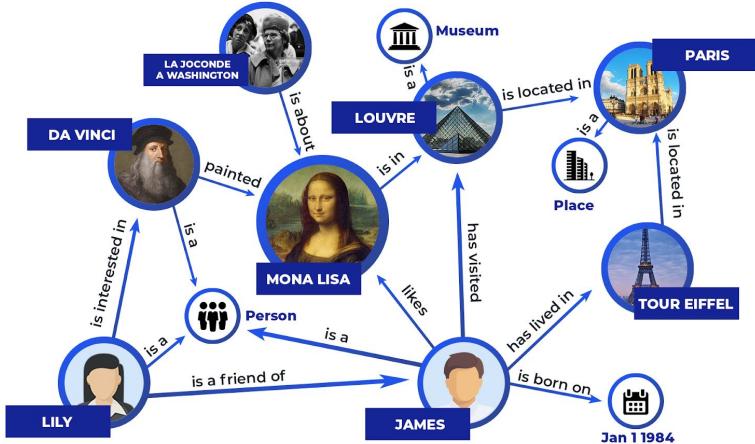
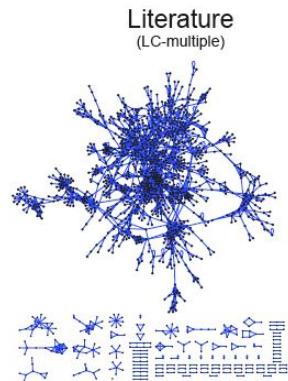
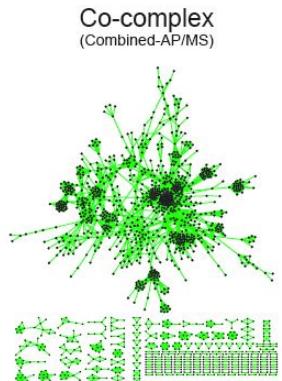
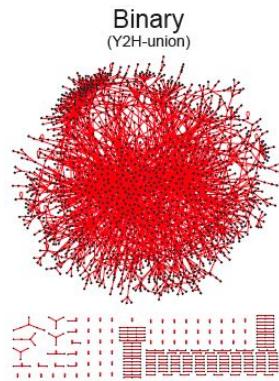
Graphs are everywhere (1/4)



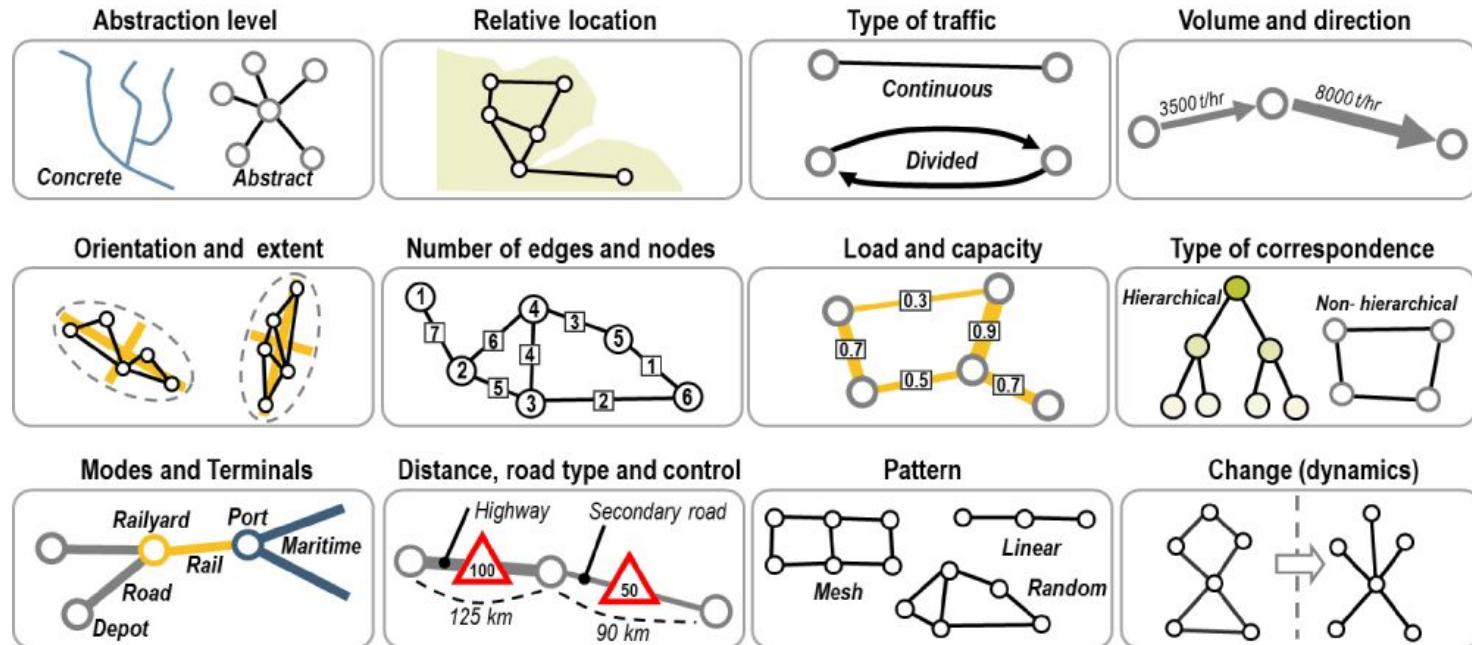
<https://linkurio.us/blog/using-neo4j-to-build-a-recommendation-engine-based-on-collaborative-filtering/>

<http://allthingsgraphed.com/2014/08/28/facebook-friends-network/>

Graphs are everywhere (2/4)

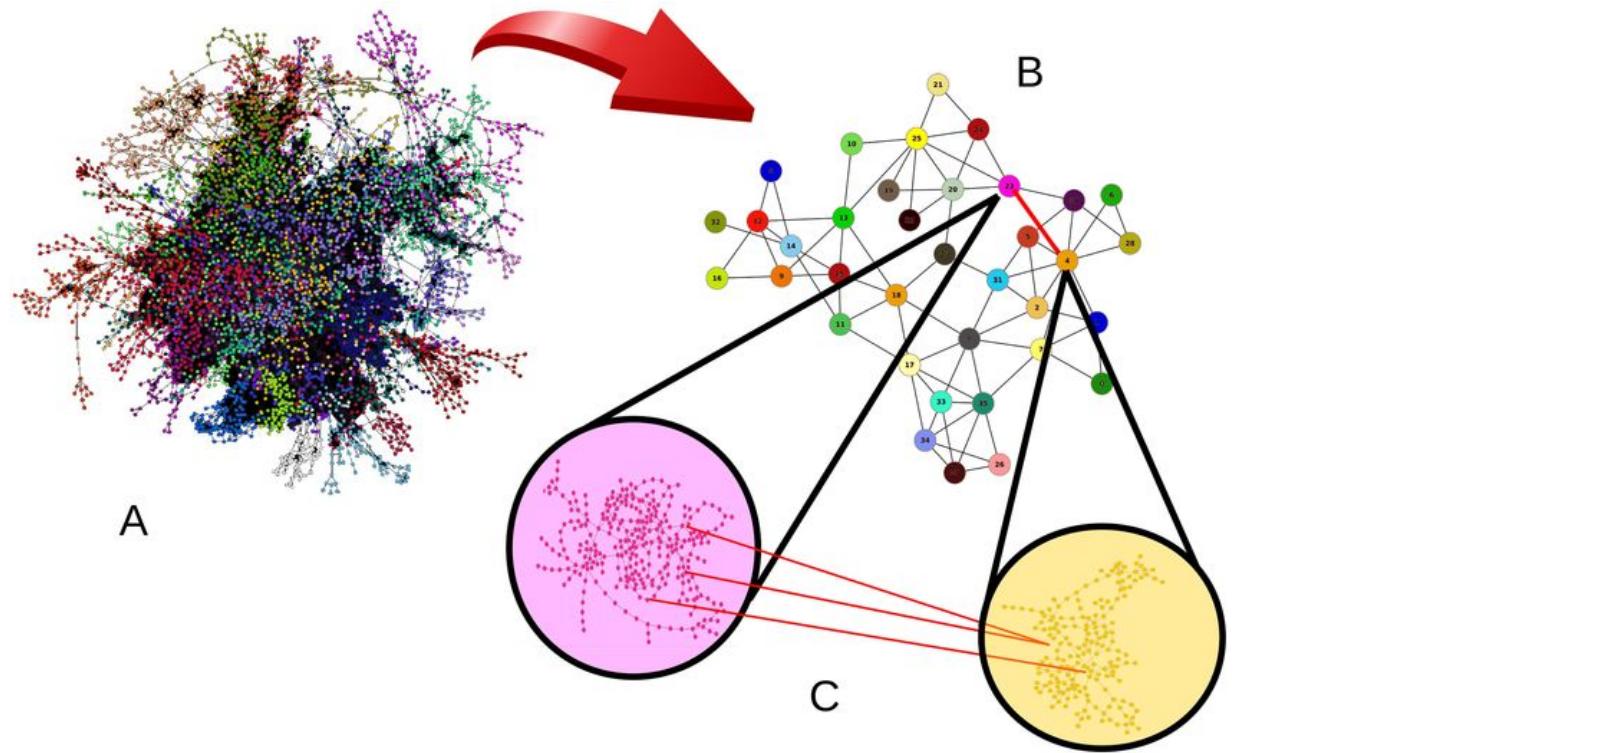


Graphs are everywhere (3/4)



https://transportgeography.org/?page_id=719

Graphs are everywhere (4/4)

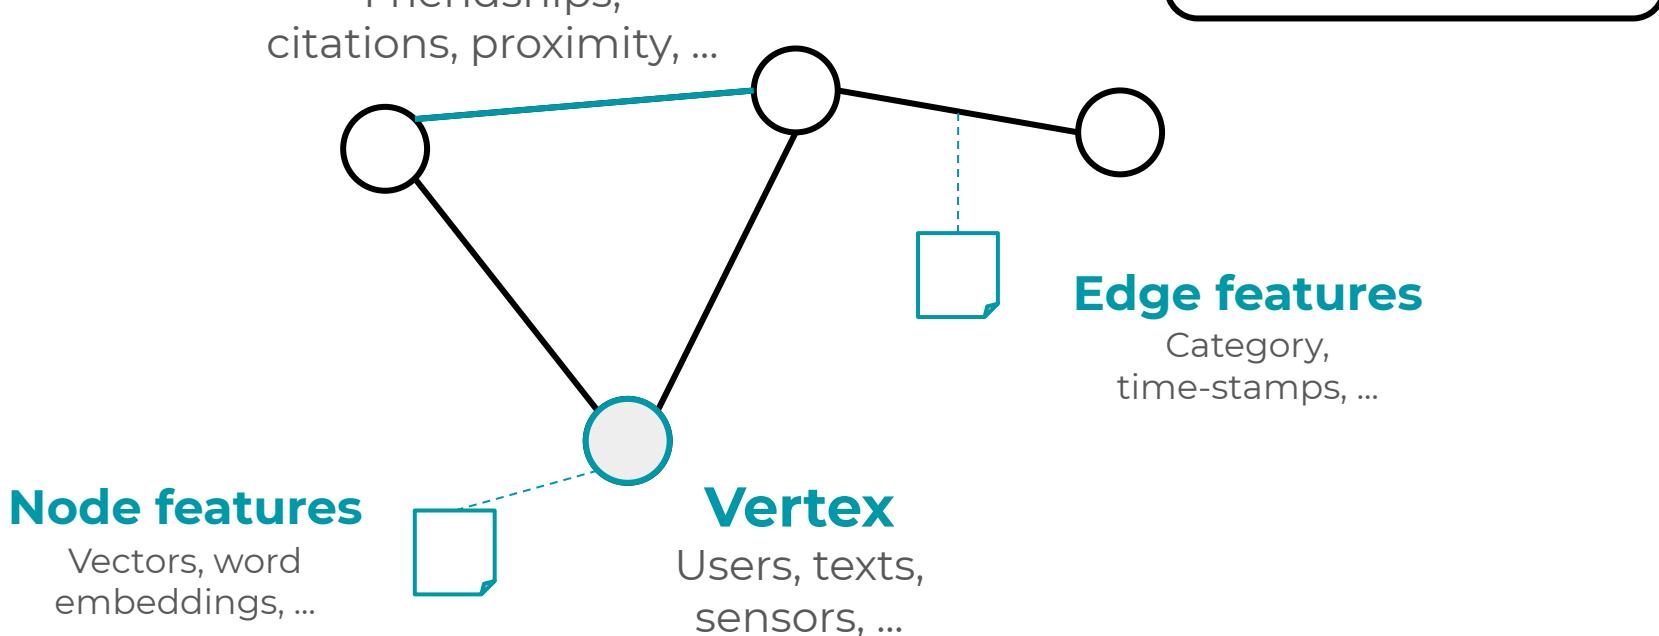


Fast Fragmentation of Networks Using Module-Based Attacks

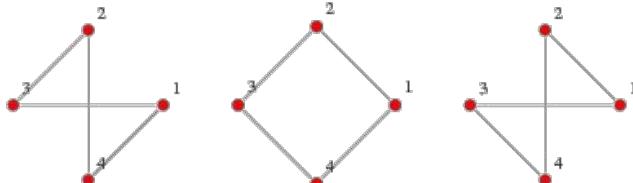
Graph and learning

Machine learning on graphs

What is a graph?



Graphs are (represented by?) matrices



$$\begin{pmatrix} 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 1 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 \end{pmatrix} \quad \begin{pmatrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 0 \end{pmatrix} \quad \begin{pmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{pmatrix}$$

Adjacency matrix

$$\mathbf{A} \in \mathbb{R}^{n \times n}$$

n vertices in the graph

Node features

$$\mathbf{X} \in \mathbb{R}^{n \times d}$$

each node has d features

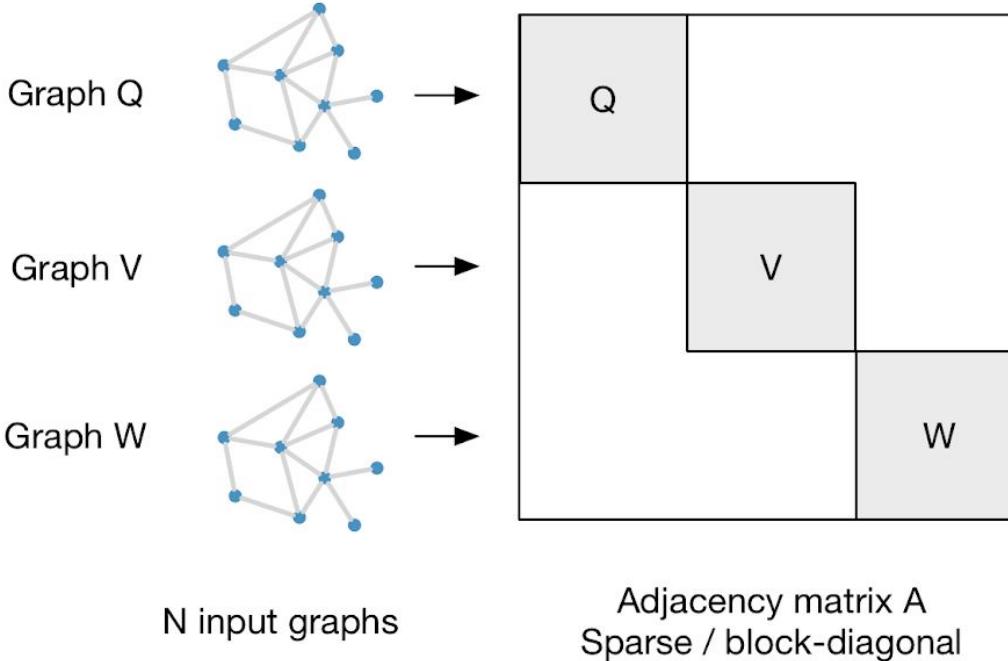
Edge features

$$\mathbf{E} \in \mathbb{R}^{e \times f}$$

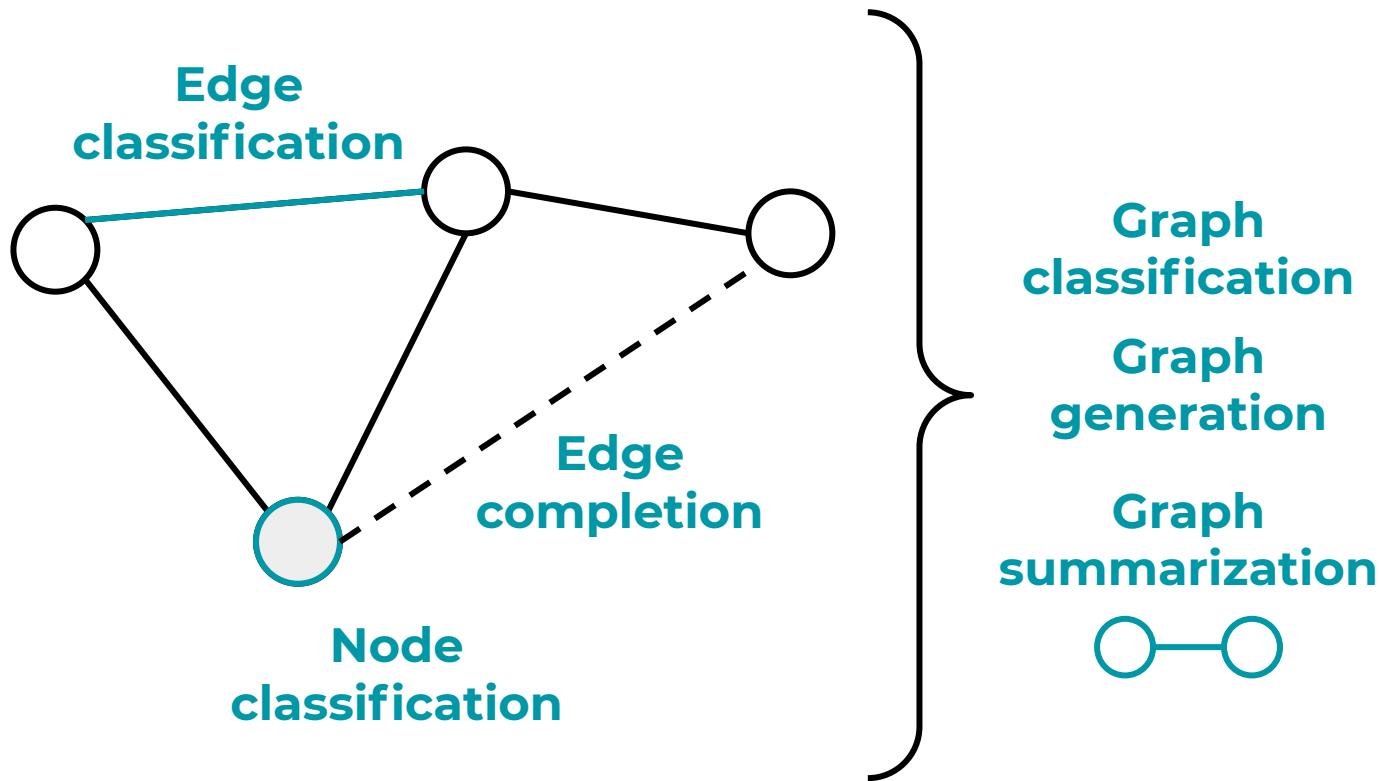
each edge has f features

[Adjacency Matrix -- from Wolfram MathWorld](#)

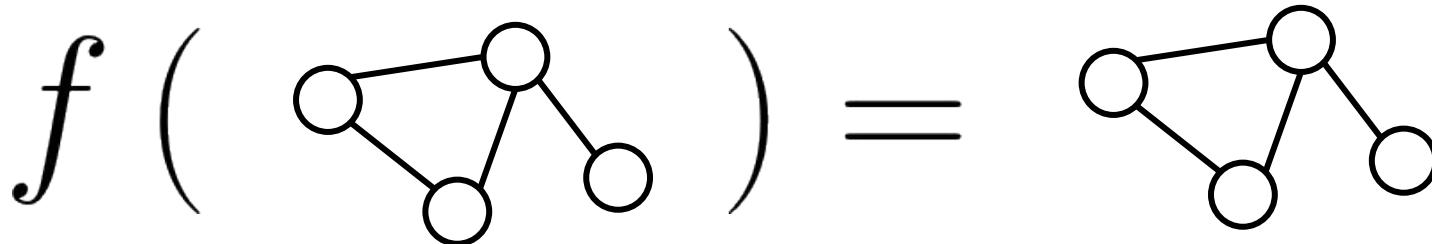
Scaling up and batching



What can we learn on a graph?



How can we learn on a graph?

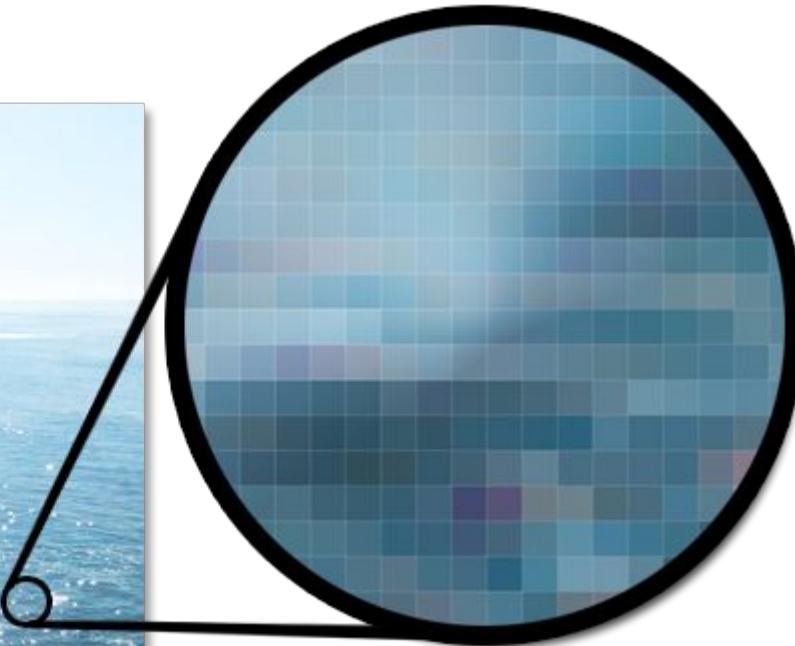
$$f \left(\begin{array}{c} \text{graph A} \end{array} \right) = \text{graph B}$$


We want to do deep learning, hence f should be
differentiable, composable, scalable.

Graph and learning

Defining a **graph layer**

Deep learning is about leveraging structure



[Generate cross-stitch patterns from any image.](#)

Image convolutions

$$f \left(\begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline & \textcolor{red}{\square} & \\ \hline & & \\ \hline & & \\ \hline \end{array} \right) = \mathbf{a}^T \textcolor{red}{\square} + \sum_{j \in \text{N}} \mathbf{c}_j^T \textcolor{blue}{\square}$$

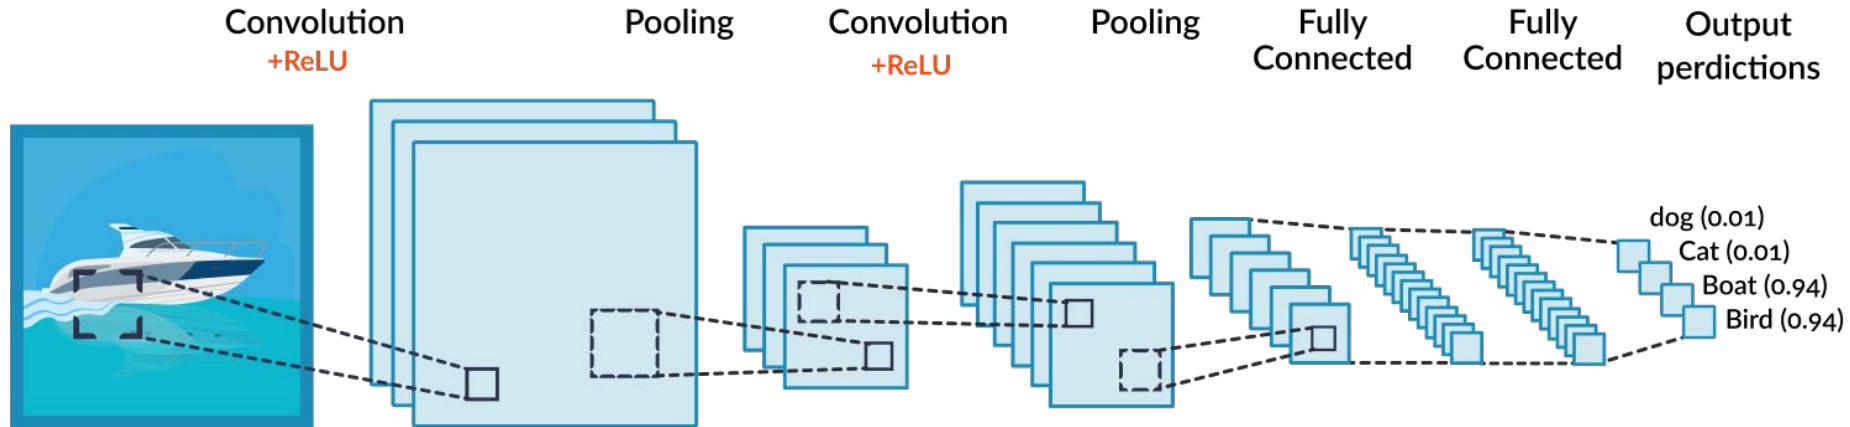
Local pixel operation

Aggregation

Neighbour pixel operation

The diagram illustrates the components of an image convolution formula. It shows a 3x3 input grid with a central red square, a weight vector \mathbf{a}^T , a bias term, and a sum over neighborhood \mathbf{c}_j^T . Arrows point from each term to its corresponding part in the formula: 'Local pixel operation' points to the central red square, 'Aggregation' points to the plus sign and the summation term, and 'Neighbour pixel operation' points to the blue square in the neighborhood set N .

Convolutional neural networks



Images vs. graphs

How much of the structure of an image do we find in graphs?

- ✓ **Locality** (neighbourhood).
- ✗ Fixed number of neighbours.
- ✗ Neighbours have a definite ordering.

Graph convolutions

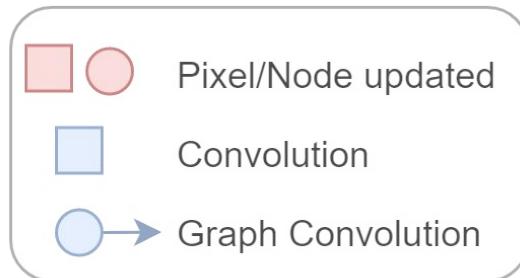
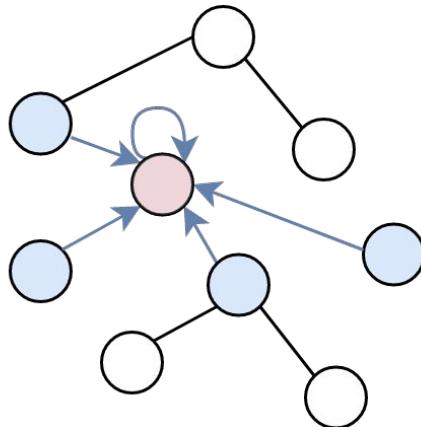
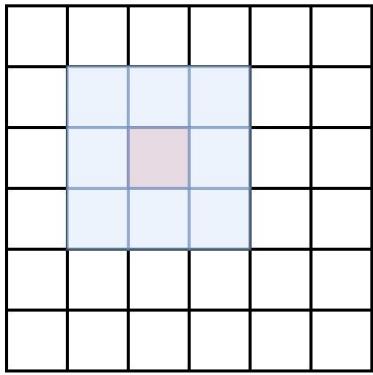
$$f \left(\begin{array}{c} \text{graph} \\ \text{(e.g., } \text{a cycle graph with one red node)} \end{array} \right) = A_{ii} \circledast \text{Local vertex operation} + \sum_{j \in \text{neighborhood}} A_{ij} \circledast \text{neighborhood graph}$$

The diagram illustrates the computation of a graph convolution. On the left, a small graph (a cycle with one red node) is shown with an arrow pointing to the term $A_{ii} \circledast \text{Local vertex operation}$. This term represents a local vertex operation where each node's feature vector is updated based on its own features and the features of its neighbors. On the right, another arrow points to the term $\sum_{j \in \text{neighborhood}} A_{ij} \circledast \text{neighborhood graph}$, which represents a global operation where each node's feature vector is updated based on the average features of all nodes in its neighborhood.

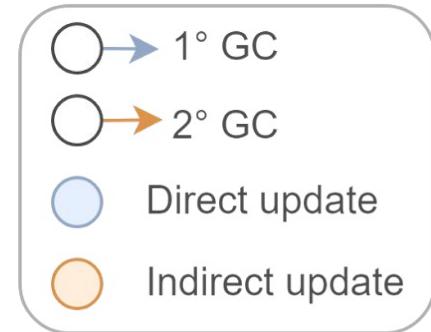
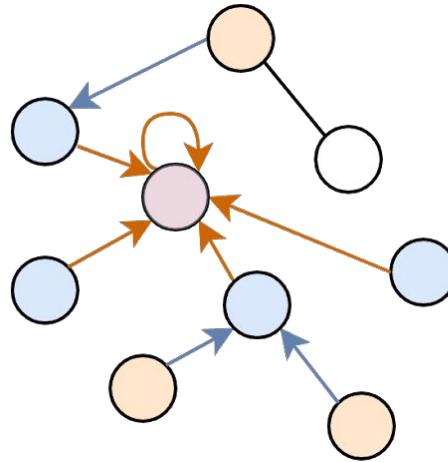
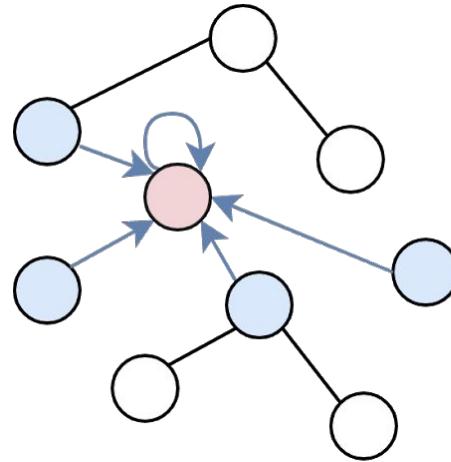
Note: the mixing coefficients are no more learnable!
Here, only local operations are trainable.

We use the
adjacency matrix
(or similar)

Graph convolution visualized



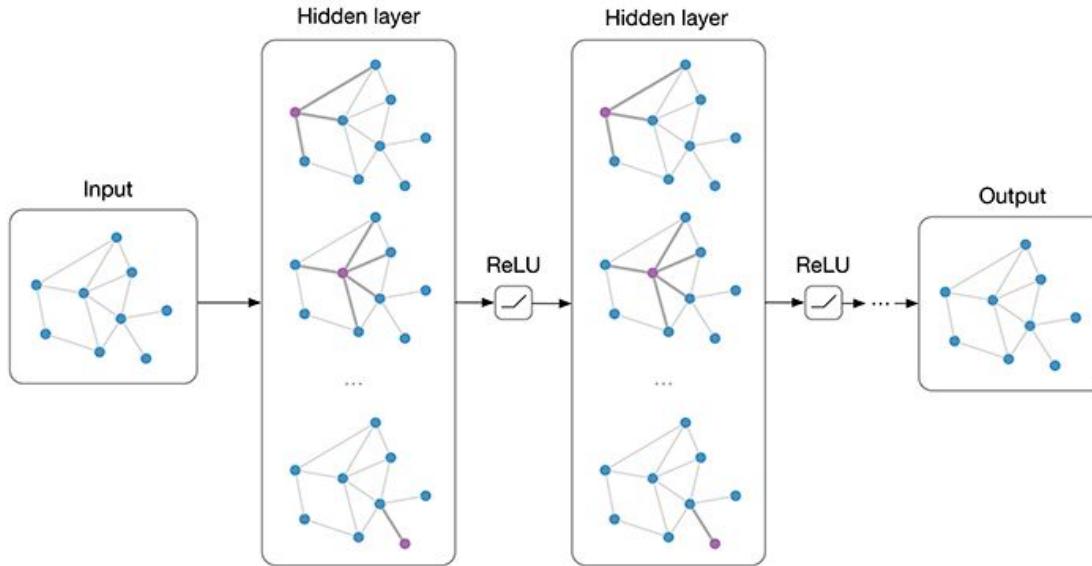
Stacking graph convolutional layers



Performing multiple updates increases the
“receptive field” of each node.

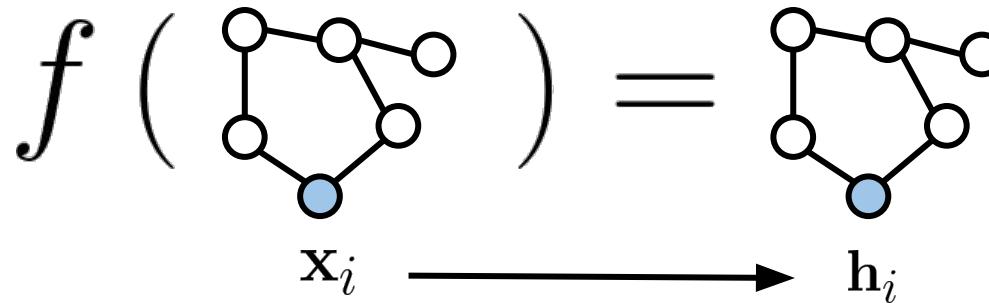
<https://spindro.github.io/post/gnn/>

Visualizing a graph convolutional network



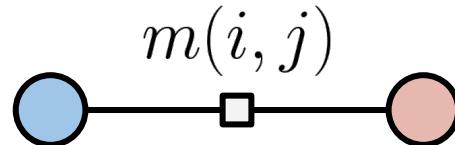
Kipf, T.N. and Welling, M., 2016. **Semi-supervised classification with graph convolutional networks**. arXiv preprint arXiv:1609.02907.

Tackling multiple tasks



1. Node classification: $\text{softmax}(\mathbf{h}_i)$
2. Edge classification: $\text{softmax}(\mathbf{h}_i^T \mathbf{h}_j)$
3. Graph classification: $\text{softmax} \left(\frac{1}{N} \sum_i \mathbf{h}_i \right)$

Message-passing neural networks

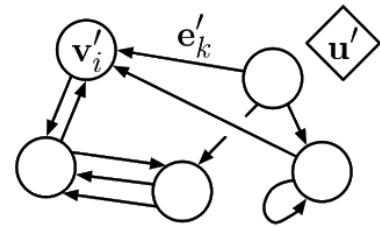
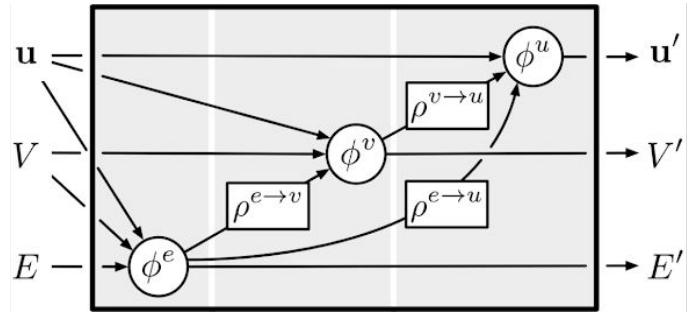
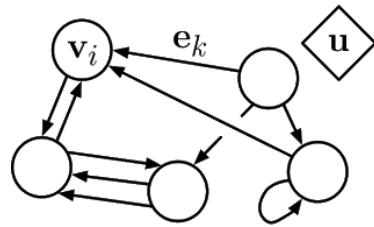


Instead of using directly the adjacency matrix, nodes can exchange **messages** with several mechanisms (e.g., attention models).



Gilmer, J., Schoenholz, S.S., Riley, P.F., Vinyals, O. and Dahl, G.E., 2017, July. **Neural message passing for quantum chemistry**. In *International Conference on Machine Learning* (pp. 1263-1272). PMLR.

A more general setup



Battaglia, P.W., Hamrick, J.B., Bapst, V., Sanchez-Gonzalez, A., Zambaldi, V., Malinowski, M., Tacchetti, A., Raposo, D., Santoro, A., Faulkner, R. and Gulcehre, C., 2018. **Relational inductive biases, deep learning, and graph networks.** arXiv preprint arXiv:1806.01261.

A zoo of techniques...

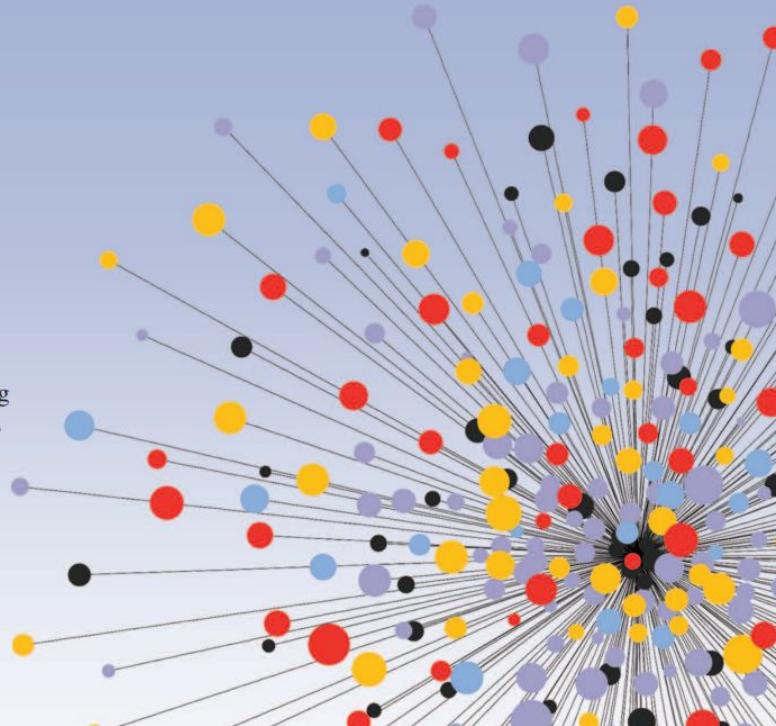
TABLE III: Summary of RecGNNs and ConvGNNs. Missing values (“-”) in pooling and readout layers indicate that the method only experiments on node-level/edge-level tasks.

Approach	Category	Inputs	Pooling	Readout	Time Complexity
GNN* (2009) [15]	RecGNN	A, X, X^e	-	a dummy super node	$O(m)$
GraphESN (2010) [16]	RecGNN	A, X	-	mean	$O(m)$
GGNN (2015) [17]	RecGNN	A, X	-	attention sum	$O(m)$
SSE (2018) [18]	RecGNN	A, X	-	-	-
Spectral CNN (2014) [19]	Spectral-based ConvGNN	A, X	spectral clustering+max pooling	max	$O(n^3)$
Henaff et al. (2015) [20]	Spectral-based ConvGNN	A, X	spectral clustering+max pooling		$O(n^3)$
ChebNet (2016) [21]	Spectral-based ConvGNN	A, X	efficient pooling	sum	$O(m)$
GCN (2017) [22]	Spectral-based ConvGNN	A, X	-	-	$O(m)$
CayleyNet (2017) [23]	Spectral-based ConvGNN	A, X	mean/graclus pooling	-	$O(m)$
AGCN (2018) [40]	Spectral-based ConvGNN	A, X	max pooling	sum	$O(n^2)$
DualGCN (2018) [41]	Spectral-based ConvGNN	A, X	-	-	$O(m)$
NN4G (2009) [24]	Spatial-based ConvGNN	A, X	-	sum/mean	$O(m)$
DCNN (2016) [25]	Spatial-based ConvGNN	A, X	-	mean	$O(n^2)$

Geometric deep learning

Michael M. Bronstein, Joan Bruna, Yann LeCun,
Arthur Szlam, and Pierre Vandergheynst

Many scientific fields study data with an underlying structure that is non-Euclidean. Some examples include social networks in computational social sciences, sensor networks in communications, functional networks in brain imaging, regulatory networks in genetics, and meshed surfaces in computer graphics. In many applications, such geometric data are large and complex (in the case of social networks, on the scale of billions) and are natural targets for machine-learning techniques. In particular, we would like to use deep neural networks, which have recently proven to be powerful tools for a broad



Bronstein, M.M. et al., 2017. **Geometric deep learning: going beyond Euclidean data.** IEEE Signal Processing Magazine, 34(4), pp.18-42.

GNNs before deep learning!

IEEE TRANSACTIONS ON NEURAL NETWORKS, VOL. 20, NO. 1, JANUARY 2009

61

The Graph Neural Network Model

Franco Scarselli, Marco Gori, *Fellow, IEEE*, Ah Chung Tsoi, Markus Hagenbuchner, *Member, IEEE*, and Gabriele Monfardini

Abstract—Many underlying relationships among data in several areas of science and engineering, e.g., computer vision, molecular chemistry, molecular biology, pattern recognition, and data mining, can be represented in terms of graphs. In this paper, we propose a new neural network model, called graph neural network (GNN) model, that extends existing neural network methods for processing the data represented in graph domains. This GNN model, which can directly process most of the practically useful types of graphs, e.g., acyclic, cyclic, directed, and undirected, implements a function $\tau(\mathbf{G}, n) \in \mathbb{R}^m$ that maps a graph \mathbf{G} and one of its nodes n into an m -dimensional Euclidean space. A supervised learning algorithm is derived to estimate the parameters of the proposed GNN model. The computational cost of the proposed algorithm is also considered. Some experimental results are shown to validate the proposed learning algorithm, and to demonstrate its generalization capabilities.

ples a function τ that maps a graph G and one of its nodes n to a vector of reals¹: $\tau(\mathbf{G}, n) \in \mathbb{R}^m$. Applications to a graphical domain can generally be divided into two broad classes, called *graph-focused* and *node-focused* applications, respectively, in this paper. In *graph-focused* applications, the function τ is independent of the node n and implements a classifier or a regressor on a graph structured data set. For example, a chemical compound can be modeled by a graph \mathbf{G} , the nodes of which stand for atoms (or chemical groups) and the edges of which represent chemical bonds [see Fig. 1(a)] linking together some of the atoms. The mapping $\tau(\mathbf{G})$ may be used to estimate the probability that the chemical compound causes a certain disease [13]. In Fig. 1(b), an image is represented by a region adjacency graph where nodes denote homogeneous regions of intensity of

Graph and learning

Some famous **applications**

Fake news detection on Twitter

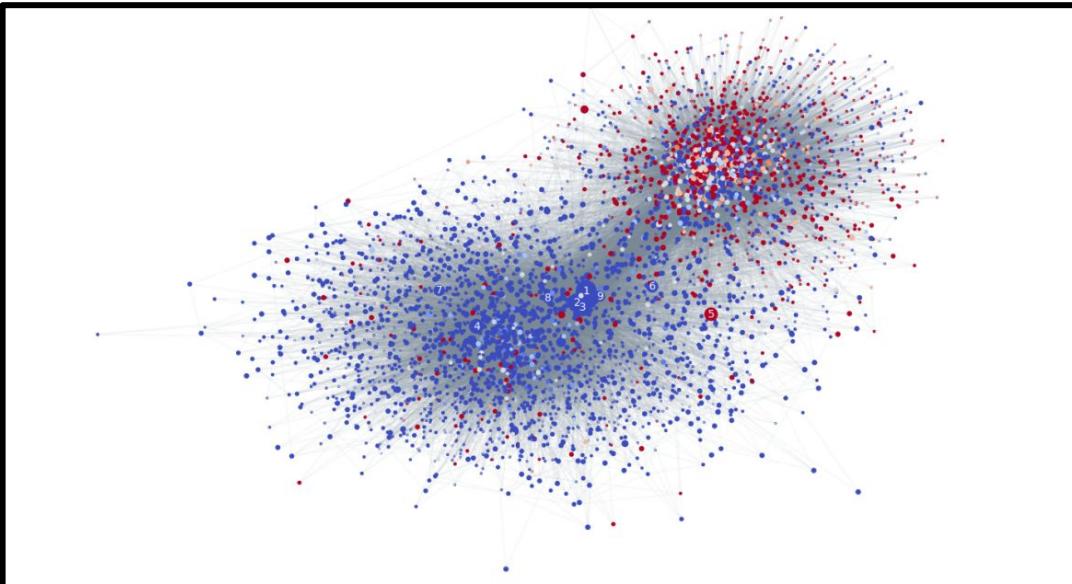


Figure 4: Subset of the Twitter network used in our study with estimated user credibility. Vertices represent users, gray edges the social connections. Vertex color and size encode the user credibility (blue = reliable, red = unreliable) and number of followers of each user, respectively. Numbers 1 to 9 represent the nine users with most followers.

Monti, F., Frasca, F., Eynard, D., Mannion, D. and Bronstein, M.M., 2019. [Fake News Detection on Social Media using Geometric Deep Learning](#). arXiv preprint arXiv:1902.06673.

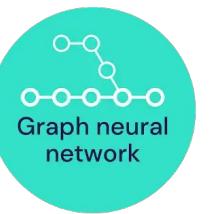
Traffic prediction on Google Maps



Analysed



Training data



Graph neural network

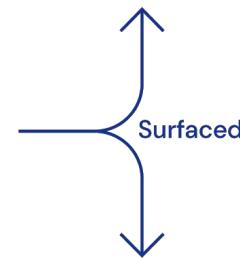
Predictions



Google Maps API



Routes ranked
by ETA



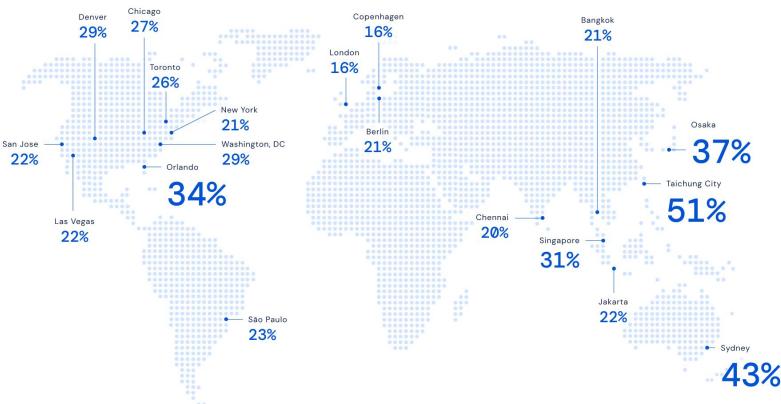
Google Maps
app



Google Maps
routing
system

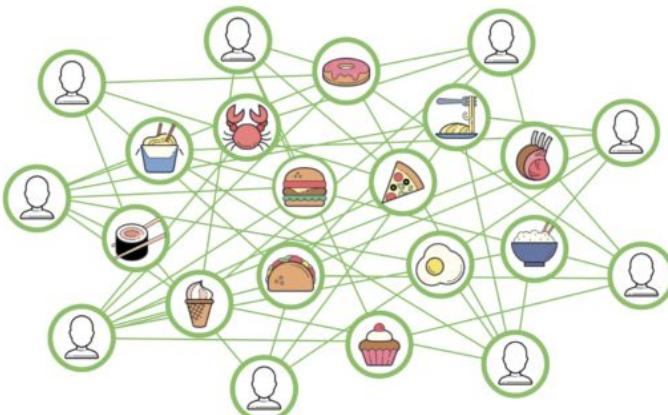
Candidate
user routes
A → B

Google Maps ETA Improvements Around the World

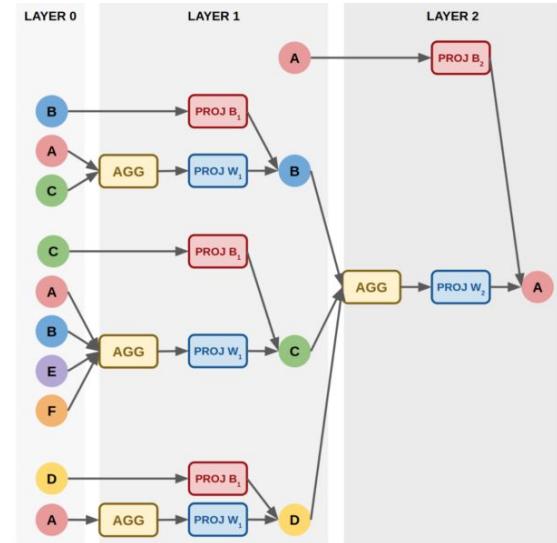
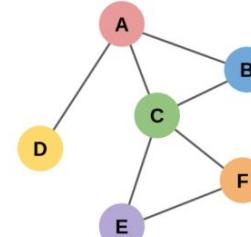


Traffic prediction with advanced Graph Neural Networks

Recommending systems in Uber

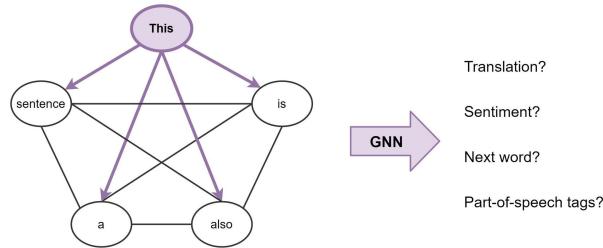


INPUT GRAPH



[Food Discovery with Uber Eats: Using Graph Learning to Power Recommendations](#)

Graph networks in classic deep learning

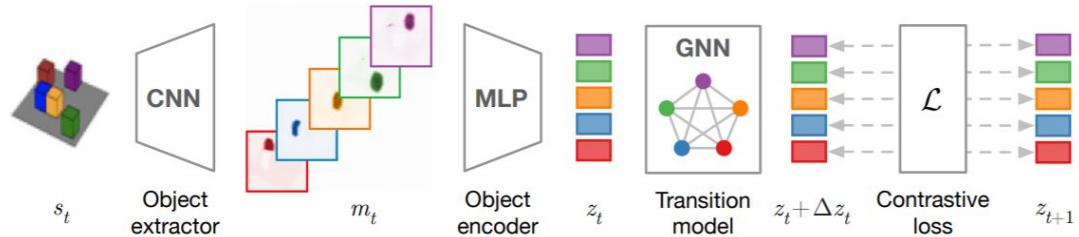


Transformers are basically GNNs on fully-connected graphs!

<https://thegradient.pub/transformers-are-graph-neural-networks/>

[\[1911.12247\] Contrastive Learning of Structured World Models](#)

GNNs can be used to include **relational reasoning** in classical models!



Graph and learning

Open problems and **challenges**

Scaling up to huge graphs

Open Graph Benchmark

Benchmark datasets, data loaders and evaluators for graph machine learning

GET STARTED

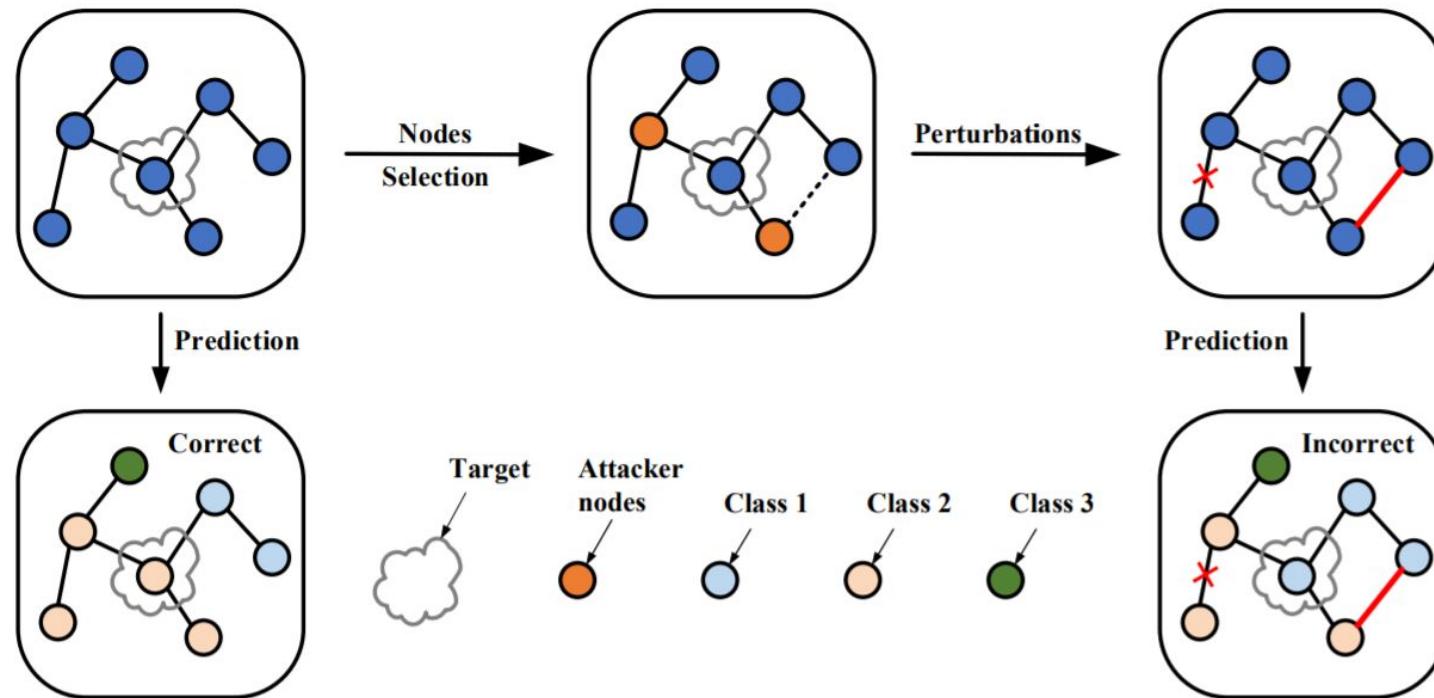
VIEW UPDATES

The Open Graph Benchmark (OGB) is a collection of realistic, large-scale, and diverse benchmark datasets for machine learning on graphs. OGB datasets are automatically downloaded, processed, and split using the [OGB Data Loader](#). The model performance can be evaluated using the [OGB Evaluator](#) in a unified manner.

OGB is a community-driven initiative in active development. We expect the benchmark datasets to

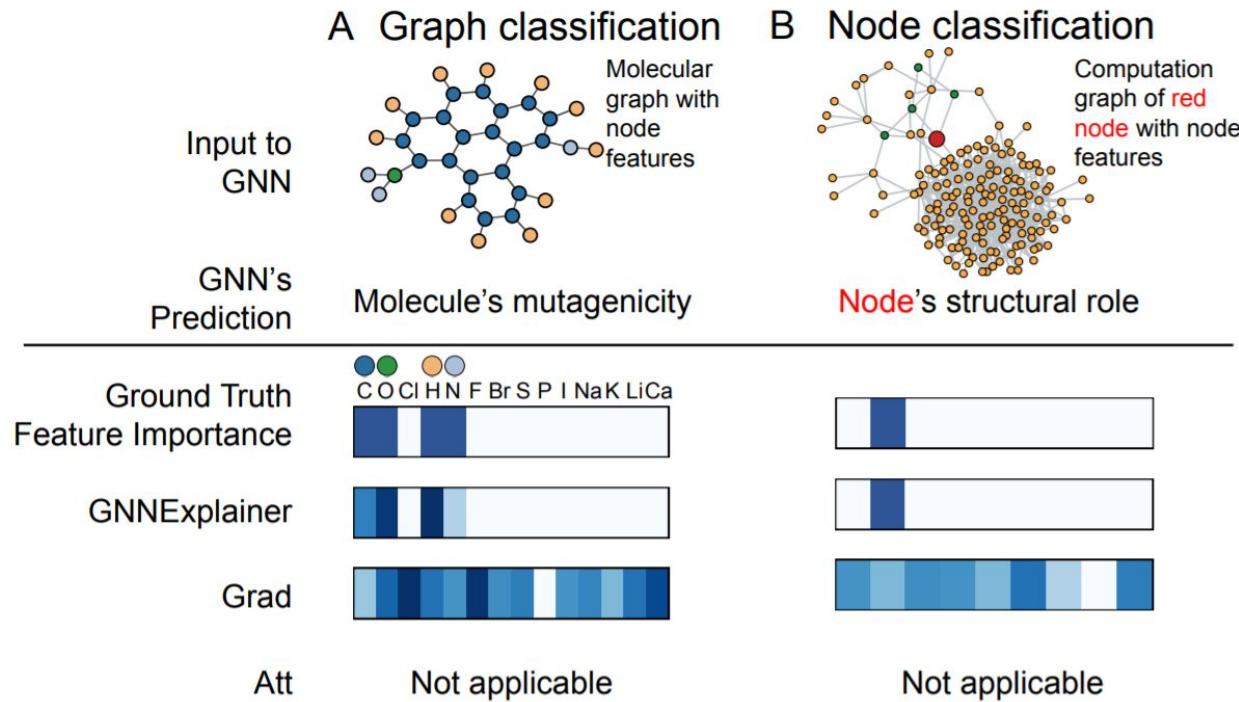


Robustness



Adversarial Attack on Large Scale Graph,
<https://github.com/safe-graph/graph-adversarial-learning-literature>

Interpretability



GNNEExplainer (Ying et al., 2019)

Software?

DeepGraphLibrary

Amazon based, production-ready



PyTorch
geometric

PyTorch, more research oriented



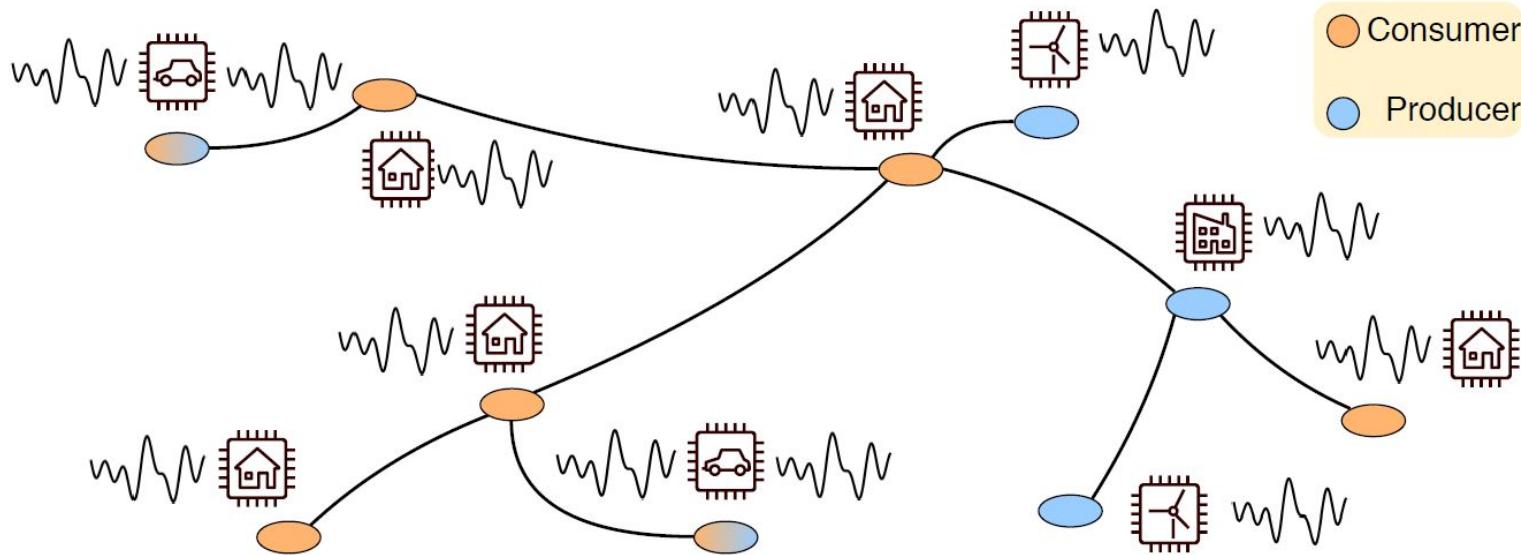
Spektral

Keras-like, TensorFlow 2.0

Distributed graphs

When graphs go **physical**

Distributed energy grids



What happens when our graph is also
physically distributed?

An example of application

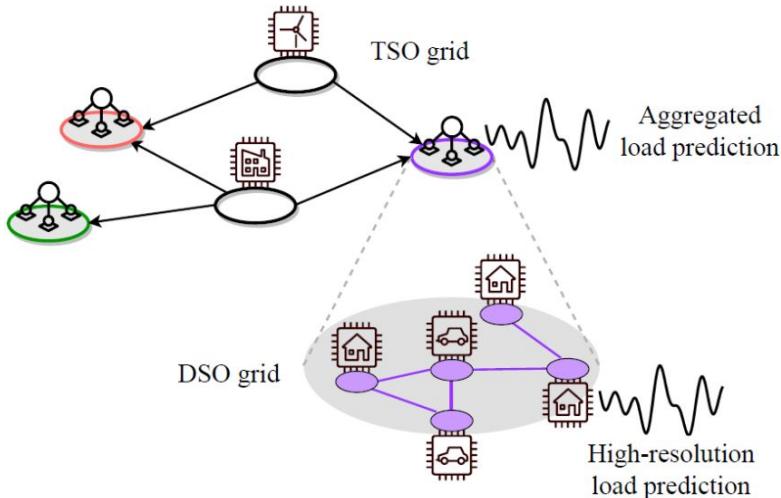
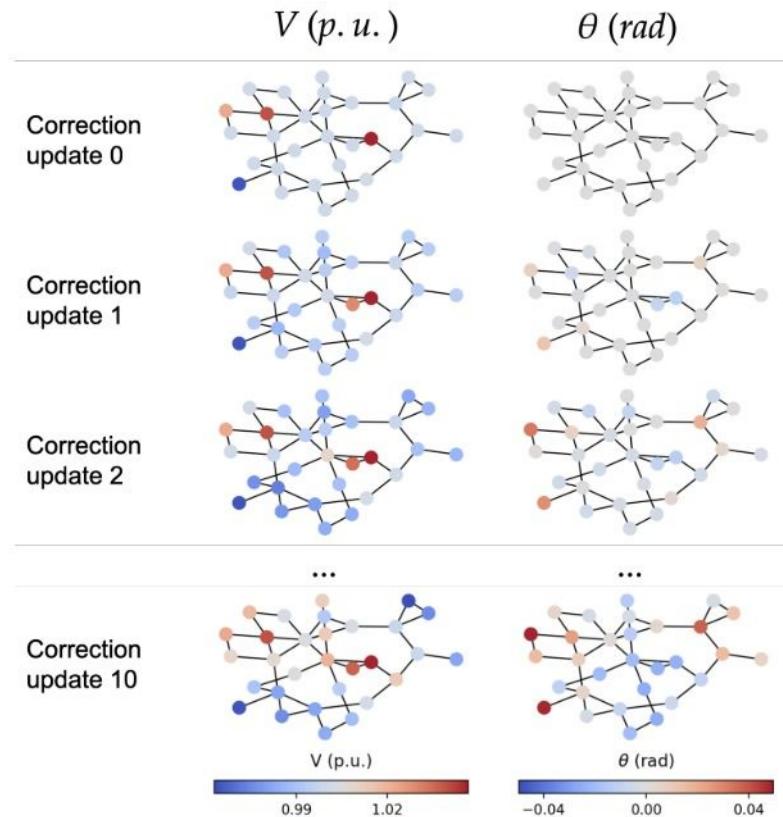
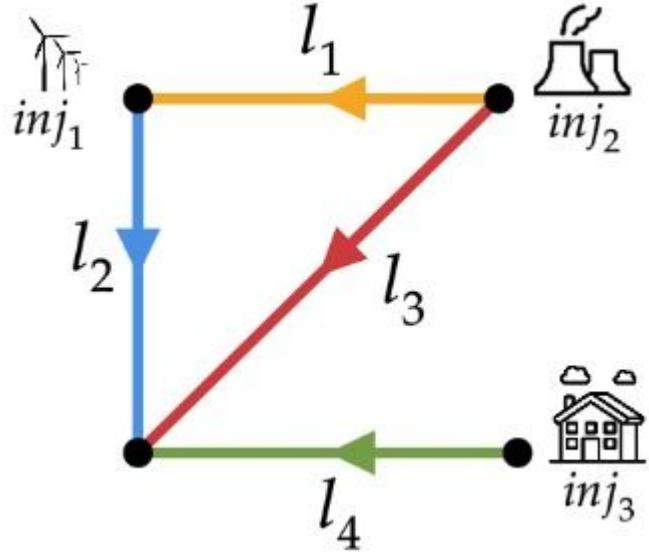


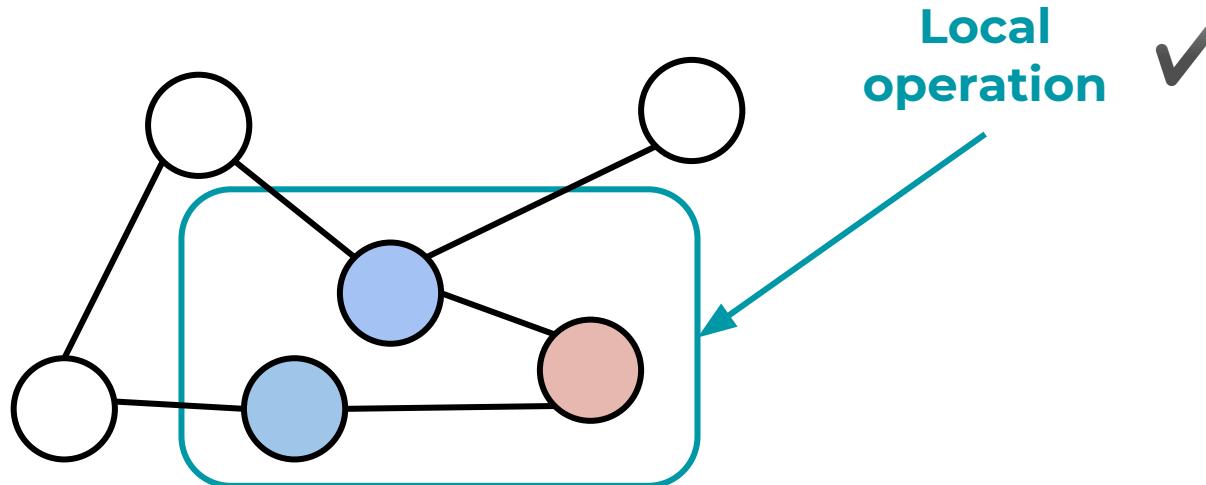
Figure 2: High-resolution load predictions are obtained on the single nodes (houses, electrical vehicles) in the low-voltage grid handled by the DSO. Through graph coarsening, the loads are aggregated at the level of resolution required to balance energy supply, provided by power plants, and the demand on the high-voltage TSO grid.

Another example

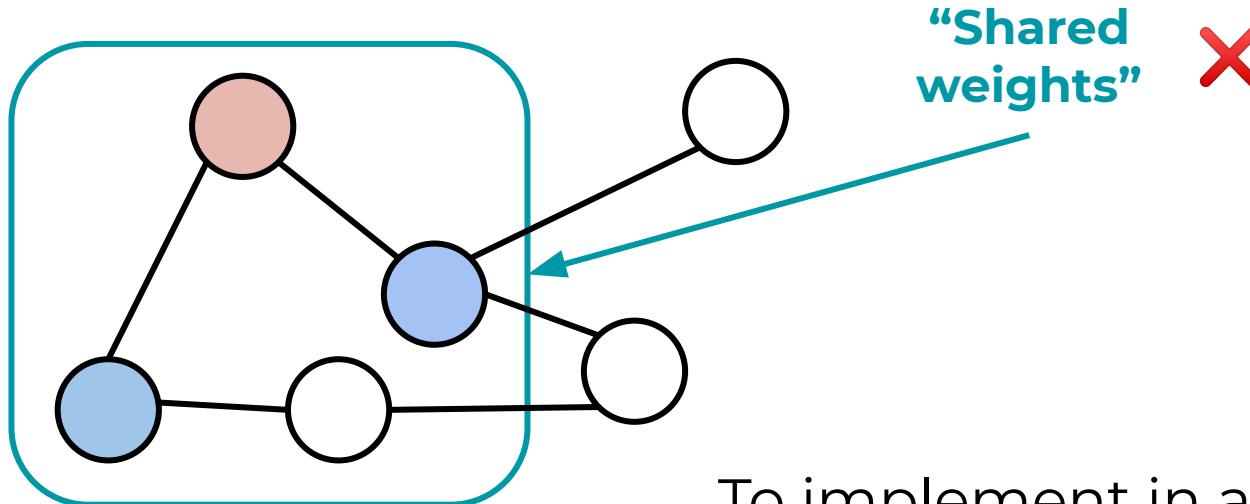


Donon, B., Donnot, B., Guyon, I. and Marot, A., 2019. **Graph neural solver for power systems**. In 2019 International Joint Conference on Neural Networks (IJCNN) (pp. 1-8). IEEE.

Partially distributed graph networks



Partially distributed graph networks



To implement in a fully decentralized fashion, nodes require perfect coordination.

Fully distributed GCNs

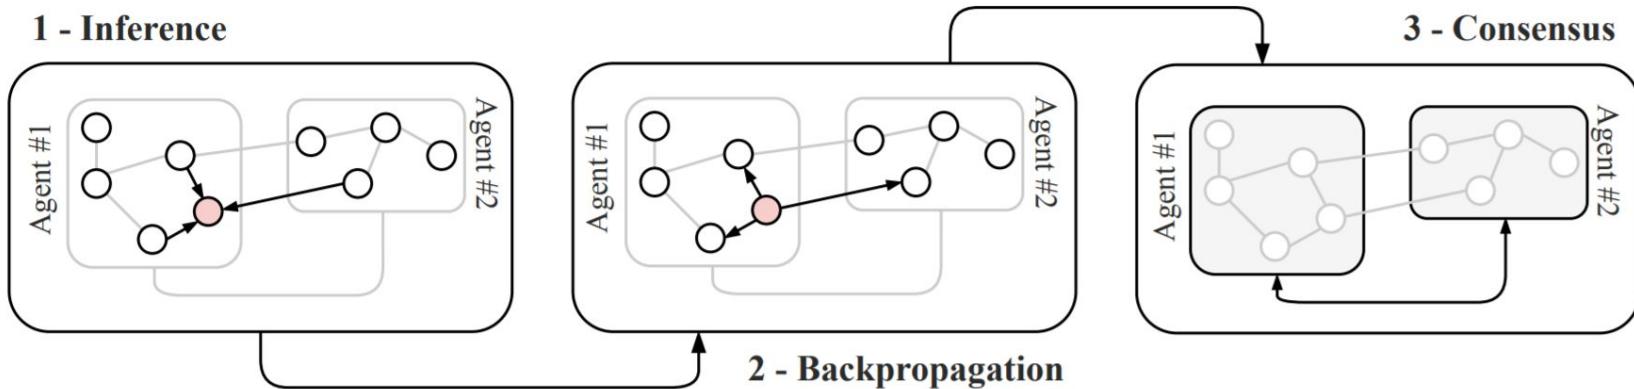


Fig. 1. Illustration of the proposed approach. In step 1, nodes communicate to perform inference. In step 2, a symmetric communication phase is executed to compute local gradients. In step 3, agents exchange local variables to asymptotically reach agreement. For steps 1-2, a representative active node is shown in red. Directed arrows show the flow of messages.

It works!

TABLE II
RESULTS OF THREE DISTRIBUTED ALGORITHMS ON THE TRAFFIC
PREDICTION BENCHMARKS FROM SECTION V-F

Dataset	Linear	GCN	Order-2 GCN
PeMSD4	0.231 ± 0.001	0.148 ± 0.002	0.126 ± 0.001
PeMSD8	0.266 ± 0.001	0.185 ± 0.002	0.147 ± 0.003

Nodes are traffic sensors in San Bernardino (California), with multiple base stations learning with no central coordination.



Conclusions

Summarizing geometric deep learning

Key takeaways

1. Graph neural networks are powerful deep learning models that are surprisingly ubiquitous.
2. Scaling up and including temporal or heterogeneous information is challenging.
3. Huge research topic, still “under the radar” in industrial contexts.

Thanks for listening!



Simone Scardapane
Assistant Professor



Indro Spinelli
PhD Student



<https://www.sscardapane.it/>

<https://spindro.github.io/>



https://twitter.com/s_scardapane

<https://twitter.com/IndroSpinelli>



BTW, check out our association:
[Italian Association for Machine Learning](#)

And why not our podcast?
[Smarter Podcast](#)